

Sonat AI

A 14-week plan to make AI the strongest thing we sell — and to be found for it

Sonat leadership · 31 July 2026 · Covering 3 August – 6 November 2026 · Decisions requested in section 10

1 The ask, in one paragraph

We have an AI answer engine that is **refusing 81% of the questions customers ask it**, and a marketing site that promises things the product does not do. We also have the one asset no competitor in this category has: **we own both the documentation tool and the answer engine**. This paper asks the board to approve fourteen weeks of work that fixes the engine, turns it into a measurable product, and makes Sonat the answer when a buyer — or a buyer's AI assistant — asks who does this. It also asks for a hard rule: **we do not market or price this until the engine is measurably fixed**.

81%

of production AI answers are refusals — 273 of 337 over 30 days

85%

of those refusals had the right documents already retrieved

\$375

our full platform price vs \$999+ that kapa.ai charges for the AI layer alone

45% vs 18%

deflection when docs are fresh vs stale — the variable we alone control

2 The strategic position

2.1 What we are actually selling

The instinct was to sell *“Enterprise gets the better AI model.”* I am recommending against that, for four reasons that survived research:

- **It is copyable in a weekend.** Every competitor buys models from the same three vendors and can switch backends by changing a string.
- **It goes stale by construction.** Our pricing page today sells “AI Answer Generator with GPT-4” — a model we have not offered for months.
- **It does not fix our actual problem.** 85% of our refusals happen *with* the right documents retrieved. That is retrieval and prompting, not model capability.
- **It invites the one question we cannot answer:** *“if it’s just GPT, why don’t I paste my docs into ChatGPT?”*

The defensible position instead: kapa.ai, Intercom Fin, Docsbot and Document360’s Eddy all answer questions over documentation they **do not own and cannot edit**. They can tell a customer their docs have a hole. **Only Sonat can fix the hole in the same tool** — because we own the authoring app, the viewer, the search index and the answer engine. And the category’s own benchmark says that loop is the variable that decides outcomes: **45% ticket deflection when the knowledge base was updated in the last 30 days, against 18% when it had not been touched in six months**. Freshness beats model. Freshness is what we control and nobody else does.

2.2 So “Sonat AI” means four things, in this order

LAYER	WHAT IT IS	WHY IT IS DEFENSIBLE
1. Answers	Cited, fast, multi-turn chat on the viewer and in the widget — and inside the customer's own Claude, ChatGPT and Cursor	Table stakes in 2026. Not having chat is the gap.

via our MCP server

2. Proof	A dashboard showing every question asked, what was answered, what failed, and the deflection value in dollars	The only AI feature that produces a number a buyer puts in a business case.
3. The loop	Every question the docs could not answer becomes a one-click draft topic in the editor	The moat. No pure-AI vendor can build it — they do not own the editor.
4. Control	Scope, per-document exclusions, audit log, provider choice, no-training guarantee, DPA	The procurement gate. Legal kills more deals in this category than features do.

Advanced model selection — which we shipped on 30 July — becomes a line item inside layer 4. It is a fine cost decision and a poor value story.

3 Where we are today — the evidence

Everything below is measured from production, not estimated. The uncomfortable parts are here because the board should see them before the plan, not after.

FINDING	DETAIL	CONFIDENCE
81% of AI answers are refusals	273 of 337 calls over 30 days returned “I don’t have enough information.” Five accounts sit at 98–100%.	reconstructed
85% of refusals had context in hand	232 of 273. The documents reached the model and it still refused. Ranking and prompting, not model quality.	reconstructed
Our own templates outrank our real content	For “how to use seo feature?”, seven of ten documents sent to the AI were placeholder text from our free manual templates. The model was handed instructions-to-an-author and correctly concluded it had nothing to answer with.	verified
One account has zero indexed content	36 of 36 customer questions returned nothing, for 30 days. Nothing alerted.	verified
We cannot measure any of this	The analytics table built for exactly this question has zero rows . Every number above was reconstructed by hand from gateway logs.	verified
Answers take 10–23 seconds	Peak 37.8s. The AI model is only ~3.2s of that; the rest is our own retrieval and cold starts. The site promises “within seconds.”	reconstructed
We only ever ask the AI the questions it is worst at	The viewer suppresses AI for queries under three words. Short queries rank beautifully; long natural-language ones are exactly what breaks the ranking. We manufactured part of our own refusal rate.	verified
The paid differentiator is a URL parameter away	Model entitlement is enforced only when saving a setting. Any anonymous caller can request our premium model for any account, billed to us at 6.7× the cost.	verified
Total AI usage is 337 calls per month	Platform-wide. Roughly 11 a day. This is the number that should worry us most — see §7.	reconstructed
Marketing is ahead of the product on every axis we tested	The site sells “GPT-4” (not shipped), “instantly” (10–23s), and that the AI “learns the content” (there is no training, no embeddings — it is keyword search plus a prompt).	verified
The site is invisible to AI assistants	Zero structured data anywhere on sonat.com, and no <code>llms.txt</code> . Nothing tells a model what we do.	verified

The consequence, stated plainly. SpecSources is already on Enterprise, already on our best model, and still refused 82.9% of questions. If we brand, price and market AI in this state, we convert a quality bug into a *paid promise* — and churn on a paid promise costs far more than churn on a free feature. **I am asking the board to hold me to a gate: no launch, no price change, no comparison marketing until refusals are measurably below 20% on three or more accounts.**

4 What the market charges — and our headroom

VENDOR	WHAT YOU GET	HEADLINE PRICE	AI BILLING
Sonat Enterprise (today)	Full platform + unlimited AI answers + unlimited viewers	\$375/mo	None — unmetered
kapa.ai Enterprise	AI answer layer only — no authoring	\$999+/mo	+ \$0.10 per answer
Document360 Enterprise	KB platform + Eddy AI suite	\$799+/project/mo	~\$0.50–1.00 per resolution
Mintlify Pro	Docs platform + in-doc AI	\$450/mo	\$0.25 per message overage
Intercom Fin	AI support agent	\$49/mo base	\$0.99 per resolution
Zendesk Advanced AI	Add-on to Zendesk	\$50/agent/mo	\$1.50–2.00 per resolution

THE WEDGE WE ARE NOT USING

A customer deflecting 500 tickets a month pays **\$500–1,000/mo for AI alone** at Intercom or Zendesk prices. We charge **\$375 for everything, with no per-answer fee and no per-viewer fee.** That headline — *unlimited AI answers, unlimited viewers, one price* — costs nothing to say and no competitor can match it.

THE PRICING HEADROOM

Once chat and analytics ship and refusals are down, Enterprise supports **\$500–650/mo** and still undercuts Document360's platform (\$799) and kapa.ai's AI layer alone (\$999). That is roughly **2× headroom** before we reach market price — but only after the gate.

5 The 14-week plan

Week 0 — 3–7 August

One day of work. Can redirect the whole quarter.

Six questions we have never answered, each cheap: how many accounts even have AI switched on; how many Business accounts have ever converted to Enterprise; what plans the worst-affected accounts are on; whether our top tier is being wrongly blocked from the models it pays for; whether anyone is using the MCP tool yet; and what an answer actually costs us. **In parallel and needing no answers: correct every false claim on the website.**

Decision meeting Friday 7 August. Go / no-go for the quarter.

Phase 1 — Weeks 1–4 — Make it work

~18 dev-days

Fix ranking so real content outranks placeholder text. Filter template stubs out of the AI's context. Rewrite the prompt so it answers partially instead of refusing wholesale. Make the search staleness check parallel and fail-open instead of silently deleting results. Reindex the broken account and alert on it happening again. Keep the answer service warm. **And restore the analytics pipeline first** — otherwise we ship blind and cannot prove any of it worked. Security fixes ride along, because the rate limiter we need for safety is the same one we will need for metering.

Exit gate: refusals under 20% on 3+ accounts · median answer under 5 seconds · no account at 100% · measured, not inferred, for 7 consecutive days. **Nothing downstream starts until this passes.**

Phase 2 — Weeks 5–8 — Make it provable

~20 dev-days

Ship **AI Insights**: every question asked, answered versus failed, which topics were cited, the ranked list of questions the documentation could not answer, and a deflection estimate in dollars. Ship the exclude-from-AI flag and the audit log. Unblock the editor AI assistant (five pull requests have been open since 23 July). Begin SOC 2 Type 2 — it is a 6–12 month clock and Type 1 is no longer accepted by enterprise buyers.

And run the cheapest possible test of this entire thesis: hand-build a *Documentation Performance Review* for five Business accounts — their traffic, their top searches, the searches that returned nothing useful, and what that cost them in support tickets. We can build these today from BigQuery with no new engineering. If nobody replies, we do not build the rest.

Exit gate: 2+ replies from the 5 review pilots · Insights live for 3 accounts.

Phase 3 — Weeks 9–12 — Make it a product, and make it visible

~33 dev-days

Sonat AI Chat on the viewer — multi-turn, cited, customer-branded, with escalation to support. **Enterprise MCP** — authenticated tokens so a customer's own Claude or ChatGPT can query their *private* manuals, with an audit trail. Something a public help centre cannot do. **The content-gap loop** — failed questions become one-click drafts. Then the launch: website, comparison pages, blog series, release notes, Product Hunt, and the visibility work in §6.

Exit gate: every public claim we make is provable on demand.

Phase 4 — Weeks 13–14 — Make it priced

Commercial only

Enterprise moves to **\$500–650**, existing accounts grandfathered 12 months. Answer allowances enforced as cost control. Upgrade offers go to the Business base — carried by their own performance review, not a feature list. Enterprise Plus stays custom, differentiated by a **CSM-led quarterly content-gap review**: a service wrapper that costs us nothing to build and cannot be commoditised.

6 Being found — by Google and by the AI assistants

Buyers increasingly do not start on Google. They ask Claude, ChatGPT or Perplexity “*what's the best AI documentation platform?*” or “*what are the alternatives to Document360?*” Today those assistants do not name us — and the reason is structural, not reputational: **there is no structured data anywhere on sonat.com, and no machine-readable description of what we do.** We are asking models to infer us from prose.

This is a first-class workstream with its own scoreboard line, in four parts:

PART	WHAT WE DO	EFFORT
Machine-readable site	Publish <code>llms.txt</code> and <code>llms-full.txt</code> . Add structured data across the site — product, pricing tiers, FAQs, how-tos, articles. Explicitly welcome the AI crawlers by name. Give every feature page a heading that is a question with a 40–60 word answer directly beneath it — that paragraph is what gets quoted. Real HTML tables, never screenshots: models cannot read pictures of comparisons.	~3 days Highest ratio on the page
Be in the data models already read	Assistants answer from aggregators, listicles and forums — not vendor sites. Submit our MCP server to the public registries (the official servers list, Smithery, mcp.so, Glama, PulseMCP, Anthropic's connector directory). Complete G2, Capterra, AlternativeTo and TrustRadius profiles. Get into the “best AI knowledge base 2026” roundups we ourselves cited in research. Product Hunt at launch. Ask our paying accounts for reviews.	Ongoing Start Phase 2 — lead times are months
Say the same thing everywhere	Models synthesize across sources; the same claim, worded identically, in many places is what makes a model confident enough to state it. One approved claim set renders into the site, the docs, every directory profile, the sales deck and the support macros. Contradictions are why a model hedges or omits you.	Process, not build
Measure it	Track referrals from chatgpt.com, perplexity.ai and claude.ai. And run a monthly LLM audit : twenty fixed questions asked of the four major assistants, logging whether Sonat appears, where, and whether the description is accurate. We baseline it in Week 5, before any of the work lands , so movement is provable rather than asserted.	Ongoing

The MCP registries are the single highest-leverage item here. We have already shipped an MCP server and our homepage already leads with it — we are genuinely early, and no other documentation platform is in those catalogues yet. Listing it puts Sonat inside the tool directories that Claude and ChatGPT users browse directly. It costs days, not weeks.

6.1 What we publish alongside it

NEW WEBSITE PAGES

- A **Sonat AI hub page** — the URL we want cited
- Feature pages: Chat, Insights, the content-gap loop, Enterprise MCP
- A **Trust page** — no-training guarantee, DPA, retention, residency, SOC 2 status. Procurement lands here.
- **Six comparison pages** — versus Document360, GitBook, Mintlify, kapa.ai, Intercom Fin, Zendesk AI
- An **interactive deflection calculator** — lead magnet and the number that sells the tier

CONTENT, ONE PIECE A WEEK FROM WEEK 4

- Ticket deflection benchmarks for 2026
- **Freshness beats model choice: 45% vs 18%** — our thesis piece, the most linkable idea we have
- What AI documentation tooling should cost in 2026
- MCP for documentation — highest citation potential
- The procurement checklist for buying AI documentation tools
- Monthly release notes, resuming in August — and a rewrite of our own product documentation, which still describes GPT-4 and GPT-3.5

Comparison pages compare **capabilities, never quality scores**. We will not publish a model-versus-model benchmark: it tells our own Business customers we knowingly ship them a worse product, and our current benchmark — six questions, one account, one run — would not survive the first prospect who asked about methodology.

7 The commercial reality, honestly

7.1 Converting the Business base has a low ceiling

	BUSINESS	ENTERPRISE	STEP
Entry price	\$40/mo	\$375/mo	9.4×
Seats included	5	15	+10 unneeded
Per user	\$8	\$25	3.1×

Thirteen active Business accounts at \$375 is **~\$4.9k MRR — the entire theoretical ceiling** of “convert our Business base.” And a five-person team cannot take that step: they must pay 9.4× and buy ten seats they will never use. **We could build every feature in this plan and convert nobody**, because the price step has no relationship to the value step.

Recommendation: let Enterprise start at five seats, around **\$125–175/mo**, decoupling capability from seat count. The trade-off is real — we would earn less from customers who would have paid \$375 — so this is **gated on one Stripe query in Week 0**. If Business → Enterprise conversions are near zero, 3× from a customer who converts beats 9.4× from one who never moves. If they are healthy, we leave pricing alone.

The larger prize is new logos and retention, not upsell. An enterprise-credible AI story is a win-rate and price-point lever on new business. And separately: **1,044 free accounts sit at 11–40 topics**, pressed against the limit and not upgrading. AI will not move them — the topic cliff will. That is a different campaign and we should not conflate the two.

7.2 The number that could invalidate all of this

337 AI answers per month, platform-wide. About eleven a day, across every account we have. There are four possible explanations and they demand four different responses:

- **AI is switched off for most accounts** → this is an onboarding and defaults problem. Cheapest win available to us.
- **It is on but nobody can find it** → a viewer UI problem, not an AI problem.
- **Users tried it, got refused, and stopped** → the engine fix is the whole job and demand follows.
- **End users do not search at all** → **the AI investment thesis is weak and we should spend this quarter elsewhere.**

Week 0 answers this with a single database query. **I want that answer before we commit the quarter**, and I am flagging to the board that it is a genuine kill criterion rather than a formality.

8 Scoreboard

MEASURE	TODAY	WEEK 4 GATE	WEEK 14 TARGET
AI refusal rate	81%	< 20% on 3+ accounts	< 15% platform-wide
Accounts refusing 98–100%	5	0	0
Median answer time	10–23s	< 5s	< 3s

Answers per month	337	1,000	5,000
Accounts with AI enabled	unknown	measured	80% of paying
Analytics pipeline	0 rows	writing	writing
Enterprise price	\$375	\$375	\$500–650
New Enterprise logos	—	—	4
Named by AI assistants (20-question audit)	baseline Wk 5	—	8 of 20
Referrals from AI assistants	~0	—	200/mo
Public reviews (G2 / Capterra)	—	—	15
SOC 2 Type 2	not started	in progress	audit window open

9 Risks

RISK	HOW IT IS MANAGED
We launch on a broken engine and turn a bug into a paid promise	The Week 4 exit gate. I am asking the board to hold it against me if I try to move it.
Week 0 reveals end users do not search at all	We stop and spend the quarter on the topic cliff and the 1,044 stuck free accounts. Real kill criterion.
We build everything and nobody converts, because of the 9.4× seat step	Five-seat Enterprise, gated on the Stripe number in Week 0.
Marketing runs ahead of the product again	One approved claim set. A claim only goes public when a named command or dashboard proves it. This is the discipline that was missing.
The 81% figure is itself wrong — it is one hand reconstruction from truncated logs	Restoring the analytics pipeline is the <i>first</i> engineering item, not a later one. The independently verified findings justify the work regardless of where the headline lands.
Competitors copy us	They can copy a model choice in a weekend. They cannot copy owning the editor. That is precisely why the pitch is the loop.
We overpromise to a Business customer who upgrades and finds the feature missing	During the pilot, analytics is sold as a manually delivered service — which is true today. Chat and the loop are sold with dates , not in the present tense. Only shipped capability is described in the present tense.
The AI-visibility work produces nothing measurable	Baselined in Week 5 and audited monthly. The machine-readable work is cheap and permanent regardless; only the outreach spend would be reconsidered.

10 Decisions requested

#	DECISION	RECOMMENDATION
1	Approve 14 weeks of engineering to fix and productise the AI answer engine	Approve ~70 dev-days across four phases

2	Adopt the Week 4 quality gate as binding — no launch, no price change, no comparison marketing until refusals are below 20% on 3+ accounts	Approve This is the item I most want on the record
3	Brand it Sonat AI , and stop selling a model number	Approve It also retires a currently false pricing-page claim
4	Correct the false marketing claims this week, independent of everything else	Approve Days of writing, no engineering
5	Publish a model-versus-model quality benchmark to justify the upgrade	Decline Compare capabilities, never scores
6	Begin SOC 2 Type 2 in Week 2	Approve 6–12 month clock; every week of delay is a deal we cannot close
7	Five-seat Enterprise at ~\$125–175/mo	Conditional On the Week 0 Stripe conversion count
8	Enterprise price increase to \$500–650	Defer To Week 13, after the gate and after chat ships
9	Fund AI-assistant visibility as a first-class workstream with its own metric	Approve Cheap, structural, and we are early on MCP
10	Accept the Week 0 kill criterion — if end users do not search, we redirect the quarter	Approve Answer due Friday 7 August

If the board approves only one thing: approve decision 2. Everything that has gone wrong here — the false claims on the site, the account that failed silently for thirty days, the refusal rate nobody noticed — came from the same root cause: **we shipped and sold faster than we measured**. The gate is the correction, and it costs nothing but patience.

Sonat · Board paper, 31 July 2026 · Prepared by Sonat leadership. Production figures marked *reconstructed* were derived by hand from API gateway logs because the analytics pipeline is not writing; restoring it is the first engineering item in Phase 1. Figures marked *verified* are independently reproducible. Competitor pricing is drawn from published pricing pages and third-party comparison sources, and should be re-checked against primary sources before external quotation. Full working plan and backlog: `.claude/docs/sonat-ai-launch-plan.md`.